

1. Project goal and original motivation

1. What concrete use case was the chatbot supposed to support? For example: answering citizen questions, helping fill forms, navigating administrative services, explaining legal/procedural requirements, etc.
2. Who was the intended user? Citizens, government employees, researchers, demo audience, or someone else?
3. Was the goal mainly a research prototype, a deployable system, a demo for EU SAI, or an evaluation benchmark?
4. What would count as “success” for this project?
5. Were there any explicit requirements from the EU SAI WP5 work package that this project was meant to satisfy?

2. Data and knowledge sources

7. What documents, websites, databases, or other sources does the chatbot use?
8. Where are these data sources stored?
9. Are the sources scraped, manually downloaded, generated, or provided by a partner?
10. What language are the documents in? Only German, or mixed German/English?
11. What type of administrative domain is covered? For example: residence permits, registration, social benefits, taxes, university administration, etc.
12. Is the current dataset complete, or was it only a small proof-of-concept dataset?
13. Are there any licensing, privacy, or usage restrictions on the data?
14. Is there any preprocessing pipeline? For example: PDF parsing, HTML cleaning, chunking, metadata extraction, deduplication.
15. What metadata is stored with each document or chunk? For example: source URL, title, date, administrative region, document type, paragraph ID.

3. RAG pipeline

16. What is the full RAG pipeline from user question to final answer?
17. How are documents chunked?
18. What embedding model is used?
19. What vector database or search backend is used?
20. Is retrieval based on dense embeddings, BM25, hybrid search, or something else?
21. How many chunks are retrieved initially?
22. Are retrieved chunks filtered before being sent to the LLM?
23. Is there a reranker? If yes, which reranker and why was it added?
24. What problem was the reranker supposed to solve?
25. Were there observed retrieval failures that motivated the current design?

- 26. Are citations or source links shown to the user?
- 27. Does the chatbot distinguish between “I found support for this answer” and “I am guessing”?

4. LLM and prompting

- 29. Which LLM is currently used?
- 30. Is the model local, API-based, open-source, or proprietary?
- 31. What prompts are used?
- 32. Is there a system prompt defining behavior, safety, language, or citation requirements?
- 33. Does the chatbot always answer in German?
- 34. Does the prompt instruct the model to refuse unsupported answers?
- 35. Were different models tested?
- 36. Why was the current model chosen?
- 37. Are there known problems with the model, such as hallucination, weak German, poor instruction following, or high latency?

5. Flow, agents, and orchestration

- 39. You mentioned “flow”? What exactly does this term mean in the current project?
- 40. Is there an agent framework being used, such as LangChain, LangGraph, LlamaIndex, Flowise, Haystack, AutoGen, or something else?
- 41. What are the steps in the current flow?
- 42. Are there multiple agents, or is it just a structured pipeline?
- 43. If there are multiple agents, what does each agent do?
- 44. Are tools being called by the model? For example: retrieval, web search, form lookup, calculator, database query.
- 45. Are agents actually necessary for the use case, or were they being explored experimentally?
- 46. What parts of the system are stable, and what parts are experimental?

6. Evaluation

- 47. What evaluation setup currently exists?
- 48. Is there a test set of questions?
- 49. If yes, where is it stored and how was it created?
- 50. Are there reference answers?
- 51. What metrics are used? For example: retrieval recall, answer correctness, faithfulness, citation accuracy, helpfulness, latency.
- 52. Are evaluations automatic, human-judged, LLM-as-judge, or mixed?
- 53. What were the latest evaluation results?
- 54. What are the biggest failure cases?

- 55. Are there example queries where the system works well?
- 56. Are there example queries where the system fails badly?
- 57. Was the purpose of evaluation to compare RAG variants, compare models, or validate the full chatbot?

7. Current implementation

- 58. Where is the code repository?
- 59. What branch or folder contains the latest working version?
- 60. How do I run the system locally?
- 61. What dependencies or environment variables are needed?
- 62. Are there API keys, model endpoints, or server credentials needed?
- 63. Is there a frontend?
- 64. Is there a backend API?
- 65. Is there a Docker setup or deployment script?
- 66. Is there a README or documentation?
- 67. Are there any known bugs or broken components?

8. Design decisions and rationale

- 71. What were the main design decisions made so far?
- 72. Why was this particular RAG setup chosen?
- 73. Why was the current reranker chosen?
- 74. Why was the current LLM chosen?
- 75. Why was the current agent or flow framework chosen?
- 76. Were there alternatives that were tried and rejected?
- 77. Were there any discussions with the Huyen about the intended direction?

9. Open problems and next steps

- 79. What were you planning to do next before the handover?
- 80. What is currently the highest-priority unfinished task?
- 81. What is the biggest technical bottleneck?

10. Minimal reproduction and examples

- 88. Could you provide one minimal working example of the system?
- 89. Could you provide 5 example user questions and the expected behavior?
- 90. Could you include screenshots or logs of the current frontend/output if available?
- 91. Could you include a short diagram of the pipeline?

Thanks a lot! Short bullet-point answers are completely fine. If a question is not yet applicable or if you don't know it or if it is not relevant just say NA or "don't know" or "not yet implemented".